

IDW seems to be a simple and easy to run method for ensuring that no one particular objective (NN, physics, etc) lags behind in improvements over the others. I am specifically interested in GradNorm, which uses the gradient magnitudes to re-weight objectives, keeping training weights equalized. It seems to be a better but more computationally intensive system.